

3. Show how to multiply the two binary numbers: 0b1010_0010 x 0b0010_0010

```
      10100010
x   00100010
-----
      00000000
     101000100
    0000000000
   00000000000
  000000000000
 1010001000000
 0000000000000
 00000000000000
-----
001010110000100
```

4. Write down the binary representation of the decimal number: 127.125 assuming IEEE 754 single precision format.

127 in binary is 1111111
0.125 in binary is 0.001

1111111.001 $\rightarrow 1.111111001 * 2^6$

Since single precision is 127, we add our exponent to it: $127 + 6 = 133 \rightarrow 10000101$

Store the sign bit $\rightarrow 0$
Store our exponent value with single precision $\rightarrow 010000101$
Store the mantissa $\rightarrow 01000010111111001$
Add zeros till 32-bit $\rightarrow 0100001011111100100000000000000$

127.125 in binary is 0100001011111100100000000000000

5. Write down the binary representation of the decimal number 1.25e10 assuming an IEEE 754 double precision format.

1.25e10 in binary is 1011101001000011101101110100000000

1011101001000011101101110100000000 $\rightarrow 1.01110100100001110110111010000000 * 2^{33}$

Since double precision is 1023, we add our exponent to it: $1023 + 33 = 1056 \rightarrow 10000100000$

Store the sign bit $\rightarrow 0$
store our exponent value with double precision $\rightarrow 010000100000$
Store the mantissa $\rightarrow 0100001000001011101001000011101101110100000000$
Add zeros till 64-bit $\rightarrow 010000100000101110100100001110110111010000000000000000000000000$

1.25e10 in binary is 010000100000101110100100001110110111010000000000000000000000000

6. Calculate, by hand, the sum of $3.95 \times 10^8 + 7.13 \times 10^2$, assuming the two numbers are stored as IEEE 754 single precision. Assume 1 guard bit, 1 round, and 1 sticky bit, and round to the nearest even number. Show all the steps.

```
3.95 * 10^8 → 395000000
7.13 * 10^2 → 713

395000000 in binary is → 10111100010110011100011000000
713 in binary is → 1011001001

10111100010110011100011000000
+ 00000000000000000000001011001001
-----
10111100010110011101110001001 → 1.0111100010110011101110001001 * 2^28

Sign bit → 0
Exponent → 127 + 28 → 155
Mantissa → 01111000101100111011100

3.95 * 10^8 + 7.13 * 10^2 = 395000713 → 01001101101111000101100111011100
```

7. Using the IEEE 754 float point single precision format, write down the pattern that would represent -0.4. Can this be represented exactly?

```
Sign Bit → 1
0 → 000000
```

```
But 0.4
0.8 ⇒ 0.4 * 2 ⇒ 0
1.6 ⇒ 0.8 * 2 ⇒ 1
1.2 ⇒ 0.6 * 2 ⇒ 1
0.4 ⇒ 0.2 * 2 ⇒ 0
0.8 ⇒ 0.4 * 2 ⇒ 0
...
```

It would get stuck in an infinite loop when trying to compute -0.4 because of the list of numbers shown above. So at some point it would need to cut off and the number represented would not be 100% accurate to the actual value.

8. The MIPS32 allows “double precision” values, but these require 64-bit registers. How did the engineers implement this on the MIPS?

Double precision values were stored into pairs which means that when storing double values in MIPS, they would need to be stored in pairs like \$f0 / \$f2 / \$f4 / \$f6 or \$f1 / \$f3 / \$f5 / \$f7 because it takes up two 32-bit registers to store the number.

9. What is the “unit of least precision”?

The number of bits in error in the least significant bits of the significand between the actual number and the number that can be represented

10. What does 0x4000 0000 represent if it's an IEEE 754 32-bit number?

0x4000 000 → 0100 0000 0000 0000 0000 0000 0000 0000

Read as an IEEE 754 value → 0 10000000 000000000000000000000000

We can assume that all zeros = 1 since the number has to equal something other than zero.

Since the maximum value for exponents is 128 for 32-bit, we store this at 2^{128} .

This makes our number 1×2^{128} or 2^{128} .

11. Why is the following C code incorrect:

```
float x = get_value();
if (x == nan) {
    printf("ERROR\n");
}
```

You cannot compare a number to a value that represents “Not a Number”

12. Define “guard digit”

Guard digit is the first of two extra bits kept on the right during intermediate calculations of floating-point numbers; They are used to improve rounding accuracy.

13. Given the 32-bit hexadecimal value: 0x8e02_0024, show what this same 32-bit value as:

a) unsigned 32-bit integer

1000 1110 0000 0010 0000 0000 0010 0100 in decimal is 2382495780

b) signed 32-bit integer

1000 1110 0000 0010 0000 0000 0010 0100 gets bit flipped to become 0111 0001 1111 1101 1111 1111 1101 1011

0111 0001 1111 1101 1111 1111 1101 1011
+ 0000 0000 0000 0000 0000 0000 0000 0001
<hr/>
0111 0001 1111 1101 1111 1111 1101 1100

0111 0001 1111 1101 1111 1111 1101 1100 in decimal is -1912471516

c) MIPS instruction

Hex → 0x8e020024 becomes 1000 1110 0000 0010 0000 0000 0010 0100 in binary.
100011 → LW
10000 → \$s0
00010 → \$v0
000000000100100 → 36

This becomes: LW \$v0 36(\$s0)

d) IEEE 754 single precision

1000 1110 0000 0010 0000 0000 0010 0100 → 1 00011100 00000100000000000100100

Sign bit → 1
Exponent → 00011100 → 28 - 127 → -99
Mantissa → 0000010000000000100100 → 131108

$00000100000000000100100 \times 2^{-99}$

$2^{(-6)} + 2^{(-18)} + 2^{(-21)} \rightarrow 0.0156292915$

Multiply by $2^{-99} \rightarrow 1.0156292915 \times 2^{-99} \rightarrow 1.60238048 \times 10^{-30}$

1000 1110 0000 0010 0000 0000 0010 0100 becomes $1.60238048 \times 10^{-30}$

14. Show the IEEE 754 encoding for:

a. 0

0 00000000 000000000000000000000000

b. -0

1 00000000 000000000000000000000000

c. +inf

0 11111111 000000000000000000000000

d. -inf

1 11111111 000000000000000000000000

e. nan

1 11111111 111111111111111111111111